

Generative Ai Fundamentals

Introduction and Applications

Part 1:

- There are two approaches to AI:
 - Discriminative AI:
 - Learns to distinguish different classes of data
 - Different, classify, identify patterns, draw conclusions
 - Generative AI:
 - Create new content based on training data
 - Starts with a prompt; output is a new text, image, audio, video, code, data
 - Both use Deep Learning
- GenAi Models:
 - Generative Adversarial Networks (GANs)
 - Variational autoencoders (VAEs)
 - Transformers
 - Diffusion Models
- Capabilities of GenAi:
 - Text generation
 - Image generation
 - Audio generation
 - Video generation
 - Code generation
 - Data generation and augmentation
 - Virtual worlds

Part 2. Foundation Models and Platforms

A. Deep Learning and Large Language Models

- Deep Learning:
 - Performed with the help of Artificial Neural Networks (ANNs):
 - ANNs are neurons

- Each neuron in the hidden layer contains inherent bias parameters
 - Connection between 2 neurons establishes weight parameters
- Organized in 3 layers:
 - Input layer
 - Hidden layers
 - Output layer
- Supervised DL:
 - Predefined output
- Unsupervised DL:
 - Applications:
 - Clustering
 - Dimensionality Reduction
- DL Neural Architecture differentiates the level of responses produced by the algo:
 - Convolutional Neural Network (CNN):
 - Conducts a convolution / mathematical operation on a previous layer
 - Able to extract useful information from images to recognize patterns, classify images and segment pics
 - Useful:
 - Image processing
 - Video recognition
 - Natural language processing
 - Recurrent Neural Network (RNN):
 - Efficient in processing sequential data such as text and speech
 - Process memory component which enables them to capture dependencies and contextual information over time
 - Useful:
 - Machine translation
 - Sentiment analysis
 - Speech recognition
 - Transformer-based models:
 - Have a 2 stack structure where there is an encoder and decoder process – an exceptionally high number of parameters

- Analyzes and captures the context and meaning of words in a hierarchical sequence and predicts the next word in the output sequence
- This leads to creation of LLM

B. Core Generative Ai Models

- There are 4 Gen Ai Models with different DL architectures and use a probabilistic approach:
 - Variational Autoencoders (VAE):
 - Work with diverse range of data – images, audio, text
 - Rapidly reduce dimensionality
 - Original Input -> Encoder -> Latent Space -> Decoder -> Reconstructed Output
 - Encoder – studies the probability distribution
 - Isolates the most useful data variables which creates the most compressed representation of the data and stores in latent space
 - Latent Space:
 - Mathematical space
 - Stores large dimensional data in a compressed format
 - Decoder:
 - Decompresses the compressed the data in the latent space to generate the desired output
 - VAE is trained in a static environment but the latent space is continuous
 - This generates new samples by random sampling
 - Produce realistic varied images
 - Usage:
 - Image synthesis: create game maps; generate Anime Avatars
 - Data compression: forecast the volatility surfaces of stocks
 - Anomaly detection: detect diseases using electrocardiogram signals
 - Generative Adversarial Networks:
 - Uses imagery and textual input
 - Uses 2 CNNs compete with each other in an adversarial game

- 1 CNN plays the role of a generator and is trained on a vast dataset to produce data samples
- Other CNN plays the role of a discriminator and tries to distinguish between real and fake samples
- Based on the discriminator's responses, the generator seeks to produce more realistic data samples
- Applications:
 - Creates new, realistic images
 - Style transfer
 - Image-to-image translation
 - Deep fakes
 - Finance – loan pricing or generating time series
 - Creates video game characters
- Challenges:
 - Difficult to train
 - Large amount of data
 - Heavy computational power
 - Can create false material
- Transformer-based Models:
 - Focuses on valuable text and filter unnecessary text
- Diffusion Models:
 - Recent addition to GenAi
 - Addresses the systematic decay of data due to noise in the latent space by applying the principles of diffusion to prevent information loss
 - Diffusion process – moves molecules from high-density to low-density similarly the diffusion models moves noise to and from the data sample
 - Step 1: Forward Diffusion: algorithm gradually add random noise to a training data
 - Step 2: Reverse Diffusion: turn the noise to around to recover the data and generate the desired output
 - Benefits:
 - Can train unlimited layers
 - Remarkable for image and video data

C. Foundation Models

- Foundation Model is a large general purpose self-supervised pre-trained on a vast amounts of unlabeled data with billions of parameters

- Large Language Models:
 - Trained on NLP
 - Develop independent reasoning which allows them to respond to queries uniquely
- Pre-Trained Models: Text-to-Text Generation
 - ML model
 - Trained on large corpus of text
 - Types of Model:
 - Statistical Model:
 - Markov Chain
 - Takes a sequence of states
 - Neural Network Models:
 - Use artificial neural networks to generate text
 - Represents complex relationships between data
 - Trained on a large corpus
 - Generates text similar to the text they are trained on
 - Use Sequence-2-sequence models or transformer models
 - Seq2Seq (Sequence to Sequence):
 - First encode the input text into a sequence of numbers then decode into a new number sequence
 - The new number sequence represents the generated text
 - Used for tasks with sequential data
 - Used:
 - Summarization
 - Speech Recognition
 - Machine Translation
 - Transformer Models:
 - Directly map the input text to the generated text

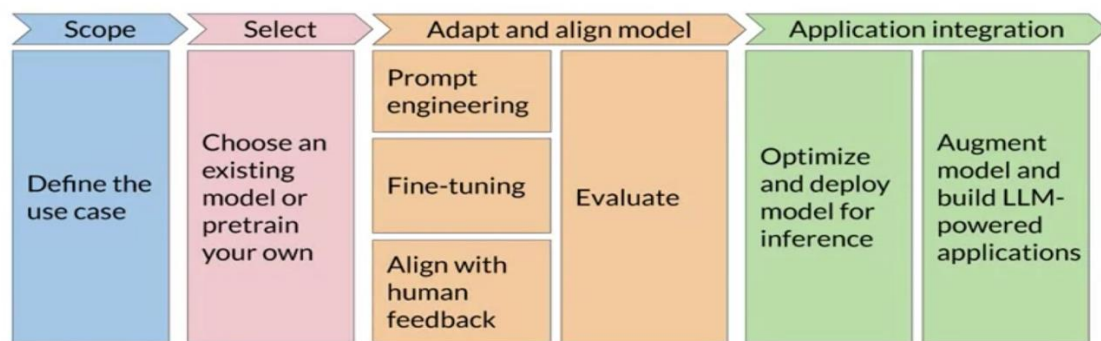
Gen Ai and Large Language Model Lifecycle

- Gen Ai is the subset of machine learning
- These models have learned to find statistical patterns in massive datasets
- Foundational / Base models:
 - GPT

- Bloom
- FLAN-T5
- PaLM
- LLaMa
- BERT

Generative AI Project Lifecycle:

Generative AI project lifecycle



1. Scope – Define the use case:
 - a. LLMs are capable of carrying out many tasks but their abilities depend strongly on the size and architecture of the model
 - b. Tasks:
 - i. Essay writing
 - ii. Summarization
 - iii. Translation
 - iv. Information retrieval
 - v. Invoke APIs and action
2. Select – Choose an existing model (foundational model) or pretrain your own:
3. Adapt and align model:
 - a. Prompt Engineering: prompt engineering is the process where you guide LLM to generate desired output.
 - i. Zero Shot
 - ii. One Shot
 - iii. Few Shot
 - b. Fine-tuning: supervised learning process
 - c. Align with human feedback

- d. Evaluate
- 4. Application Integration:
 - a. Optimize and deploy model for inference
 - b. Augment model and build LLM powered applications

